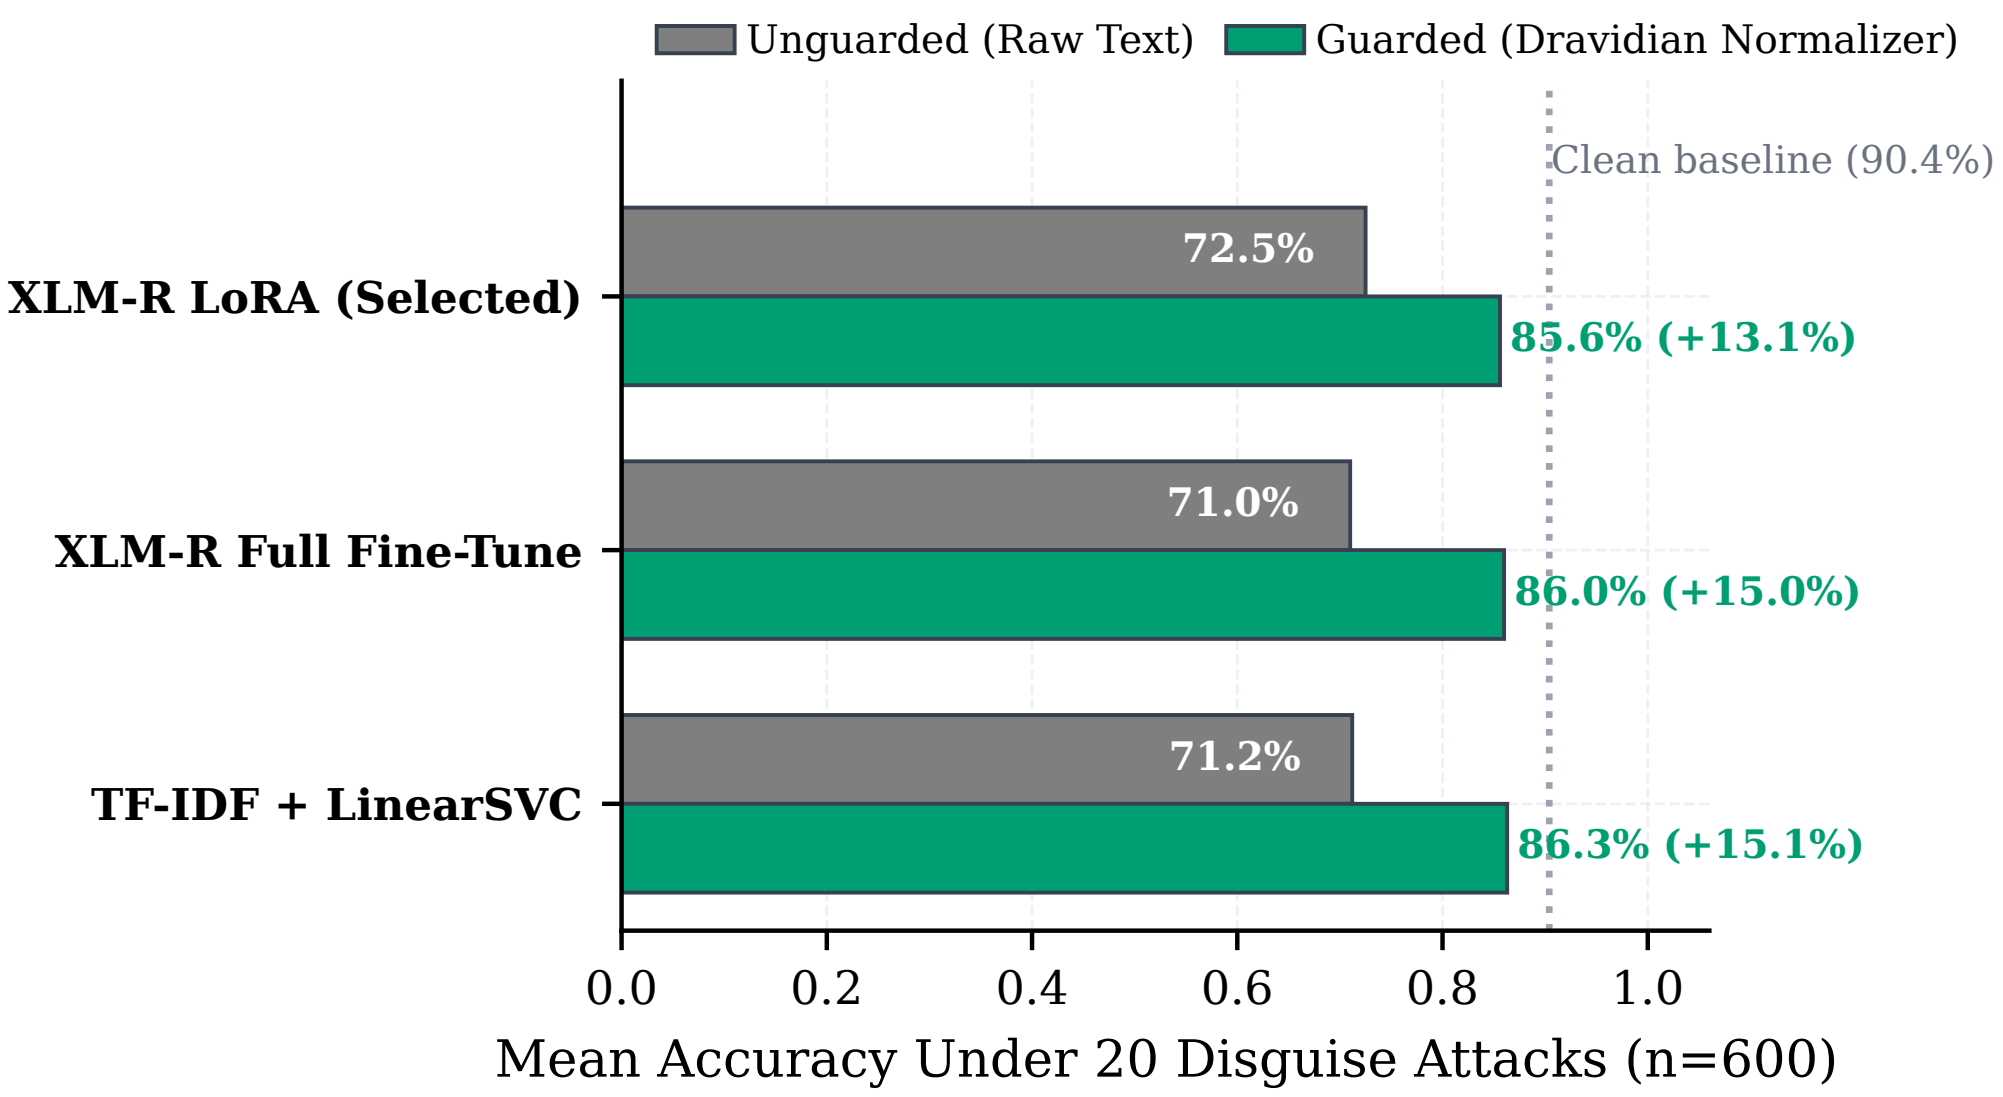
Unguarded (Raw Text)
Guarded (Dravidian Normalizer)
Clean baseline (90.4%)
XLM-R LoRA (Selected)
72.5%
85.6% (+13.1%)
XLM-R Full Fine-Tune
71.0%
86.0% (+15.0%)
TF-IDF + LinearSVC
71.2%
86.3% (+15.1%)
0.0
0.2
0.4
0.6
0.8
1.0
Mean Accuracy Under 20 Disguise Attacks (n=600)